

Assignment Project Exam Help

I <https://eduassistpro.github.io/> re

COMSC 2

Add WeChat edu_assist_pro

Outline

- Intel Microprocessors
 - IA-32 Registers
 - Instruction Execution
 - IA-32 Memory Manag
- Assignment Project Exam Help
<https://eduassistpro.github.io/>
Add WeChat edu_assist_pro

Intel Microprocessors

- Intel introduced the 8086 microprocessor in 1979
- 8086, 8087, 8088, and 80186 processors
 - 16-bit processors with 16-bit data bus
 - 16-bit data bus and 20-bit address bus
 - Physical address space = 2^{20} bytes = 1 MB
 - 8087 Floating-Point co-processor
 - Uses segmentation and real-address mode to address memory
 - Each segment can address 2^{16} bytes = 64 KB
 - 8088 is a less expensive version of 8086
 - Uses an 8-bit data bus
 - 80186 is a faster version of 8086

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Intel 80286 and 80386 Processors

- 80286 was introduced in 1982
 - 24-bit address bus $\Rightarrow 2^{24}$ bytes = 16 MB address space
 - Introduced **protected mode**
 - Segmentation in protected mode
- 80386 was introduced in 1985
 - First **32-bit processor** with 32-bit general purpose registers
 - First processor to define the IA-32 architecture
 - 32-bit data bus and 32-bit address bus
 - 2^{32} bytes \Rightarrow 4 GB address space
 - Introduced **paging, virtual memory**, and the **flat memory model**
 - Segmentation can be turned off

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Intel 80486 and Pentium Processors

- 80486 was introduced 1989
 - Improved version of Intel 80386
 - On-chip **Floating-Point unit** (DX versions)
 - On-chip unified **Instruction/**
 - Uses **Pipelining**: can execute up to 1 instruction
- Pentium (80586) was introduced in 1993
 - Wider 64-bit data bus, but address bus is still 32 bits
 - Two execution pipelines: U-pipe and V-pipe
 - **Superscalar** performance: can execute 2 instructions per clock cycle
 - Separate 8 KB instruction and 8 KB data caches
 - **MMX instructions** (later models) for multimedia applications

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Intel P6 Processor Family

- P6 Processor Family: Pentium Pro, Pentium II and III
- Pentium Pro was introduced in 1995
 - **Three-way superscalar**: can execute 3 instructions per clock cycle
 - 36-bit address bus \Rightarrow up to 64 GB address space
 - Introduced dynamic execution
 - **Out-of-order** and **speculative** execution
 - Integrates a 256 KB second level **L2 cache** on-chip
- Pentium II was introduced in 1997
 - Added **MMX instructions** (already introduced on Pentium MMX)
- Pentium III was introduced in 1999
 - Added **SSE instructions** and eight new 128-bit XMM registers

Pentium 4 and Xeon Family

- Pentium 4 is a seventh-generation x86 architecture
 - Introduced in 2000
 - New micro-architecture design called Intel Netburst
 - Very deep instruction pipeline, allowing for very high frequencies
 - Introduced the SSE (Superscalar Streamer Extension) to SSE
 - Tuned for multimedia and operating system M registers
- In 2002, Intel introduced Hyper-Threading technology
 - Allowed 2 programs to run simultaneously, sharing resources
- Xeon is Intel's name for its server-class microprocessors
 - Xeon chips generally have more cache
 - Support larger multiprocessor configurations

Pentium-M and EM64T

- Pentium M (**Mobile**) was introduced in 2003
 - Designed for **low-power** laptop computers
 - Modified version of Pentium III, optimized for power efficiency
 - Large second-level cache (2 M
 - Runs at lower clock than Pentium III
- Extended Memory 64-bit Technology (EM64T)
 - Introduced in 2004
 - 64-bit superset of the IA-32 processor architecture
 - 64-bit general-purpose registers and integer support
 - Number of general-purpose registers increased from 8 to 16
 - 64-bit pointers and flat virtual address space
 - Large physical address space: up to $2^{40} = 1$ Terabytes

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Intel Core Microarchitecture

- 64-bit cores
- Wide dynamic execution (execute four instructions simultaneously)
- Intelligent power cap
- Advanced smart cache
- Smart memory access (memory disambiguation)
- Advanced digital media boost

See the demo at

http://www.intel.com/technology/architecture/coremicro/demo/demo.htm?iid=tech_core+demo

CISC and RISC

- CISC – Complex Instruction Set Computer
 - Large and complex instruction set
 - Variable width instructions
 - Requires microcode interpreter
 - Each instruction is broken down into multiple micro-operations
 - Example: Intel
- RISC – Reduced Instruction Set Computer
 - Small and simple instruction set
 - All instructions have the same width
 - Simpler instruction formats and addressing modes
 - Decoded and executed directly by hardware
 - Examples: ARM, MIPS, PowerPC, SPARC, etc.

Next ...

- Intel Microprocessors

- IA-32 Registers Assignment Project Exam Help

- Instruction Execution

- IA-32 Memory Manag

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Basic Program Execution Registers


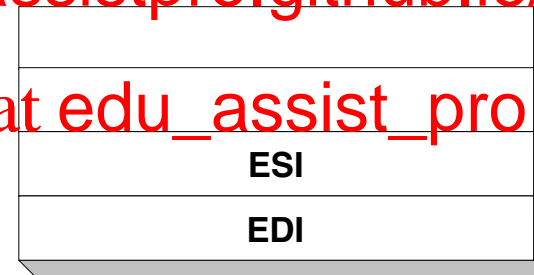
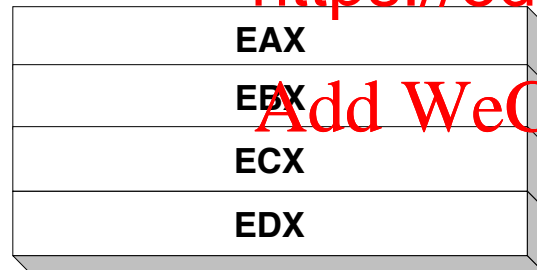
- Registers are high speed memory inside the CPU
 - Eight 32-bit general-purpose registers
 - Six 16-bit segment registers
 - Processor Status Flags

Assignment Project Exam Help


Pointer (EIP)

³ <https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



16-bit Segment Registers



General-Purpose Registers

- Used primarily for arithmetic and data movement
 - `mov eax, 10` move constant 10 into register eax

- Specialized uses of Registers

- EAX – **Accumulator** register
 - Automatically used by multiplication and division instructions
- ECX – **Counter** register
 - Automatically used by LOOP instructions
- ESP – **Stack Pointer** register
 - Used by PUSH and POP instructions, points to top of stack
- ESI and EDI – **Source Index** and **Destination Index** register
 - Used by string instructions
- EBP – **Base Pointer** register
 - Used to reference parameters and local variables on the stack


Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Accessing Parts of Registers

- EAX, EBX, ECX, and EDX are 32-bit **Extended** registers
 - Programmers can access their 16-bit and 8-bit parts
 - Lower 16-bit of EAX is named AX
 - AX is further divided into
 - AL = lower 8 bits
 - AH = upper 8
- ESI, EDI, EBP, ESP have 16-bit names for lower half



32-bit	16-bit	8-bit (high)	8-bit (low)
EAX	AX	AH	AL
EBX	BX	BH	BL
ECX	CX	CH	CL
EDX	DX	DH	DL

32-bit	16-bit
ESI	SI
EDI	DI
EBP	BP
ESP	SP

Accessing Parts of Registers

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Special-Purpose & Segment Registers

- EIP = Extended Instruction Pointer
 - Contains address of next instruction to be executed
- EFLAGS = Extended Flags Register
 - Contains status and control flags
 - Each flag is a bit
- Six 16-bit Segment Registers
 - Support segmented memory
 - Six segments accessible at a time
 - Segments contain distinct contents
 - Code
 - Data
 - Stack

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

EFLAGS Register

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

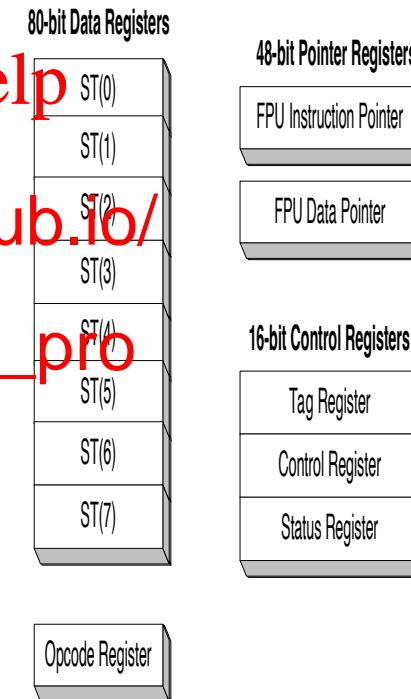
- **Status Flags**
 - Status of arithmetic and logical operations
- **Control and System flags**
 - Control the CPU operation
- Programs can set and clear individual bits in the EFLAGS register

Status Flags

- Carry Flag
 - Set when **unsigned** arithmetic result is out of range
- Overflow Flag
 - Set when **signed** arithmetic result is out of range
- Sign Flag
 - Copy of **sign bit** <https://eduassistpro.github.io/>
- Zero Flag
 - Set when result is **zero**
- Auxiliary Carry Flag
 - Set when there is a **carry from bit 3 to bit 4**
- Parity Flag
 - Set when parity is **even**
 - Least-significant **byte** in result contains **even number of 1s**

Floating-Point, MMX, XMM Registers

- Floating-point unit performs high speed FP operations
- Eight 80-bit floating-point data registers
 - ST(0), ST(1), ..., ST(7)
 - Arranged as a stack
 - Used for floating-point arithmetic
- Eight 64-bit MMX registers
 - Used with MMX instructions
- Eight 128-bit XMM registers
 - Used with SSE instructions



Next ...

- Intel Microprocessors
- IA-32 Registers
- Instruction Execution
- IA-32 Memory Manag

Assignment Project Exam Help

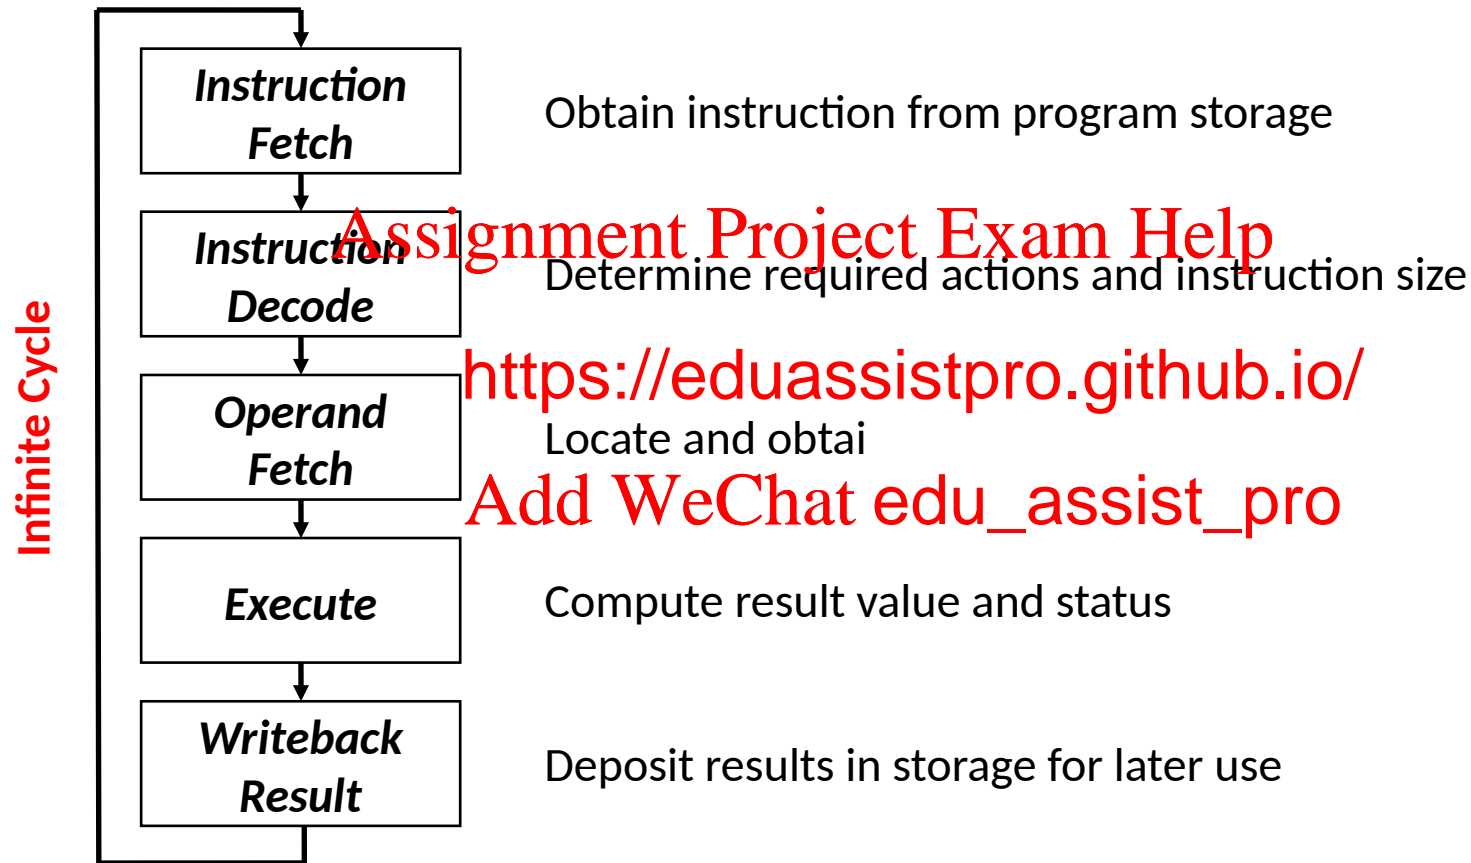
<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Fetch-Execute Cycle

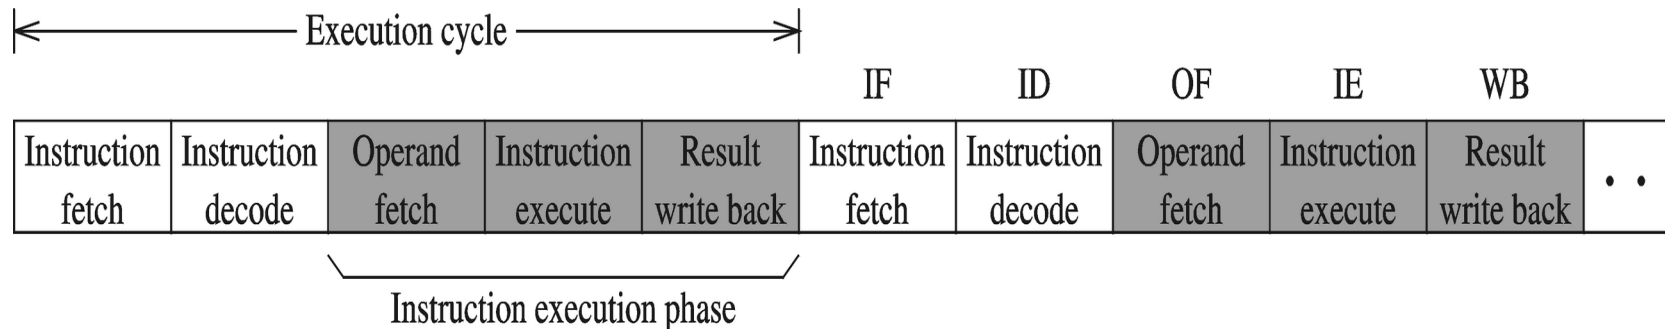
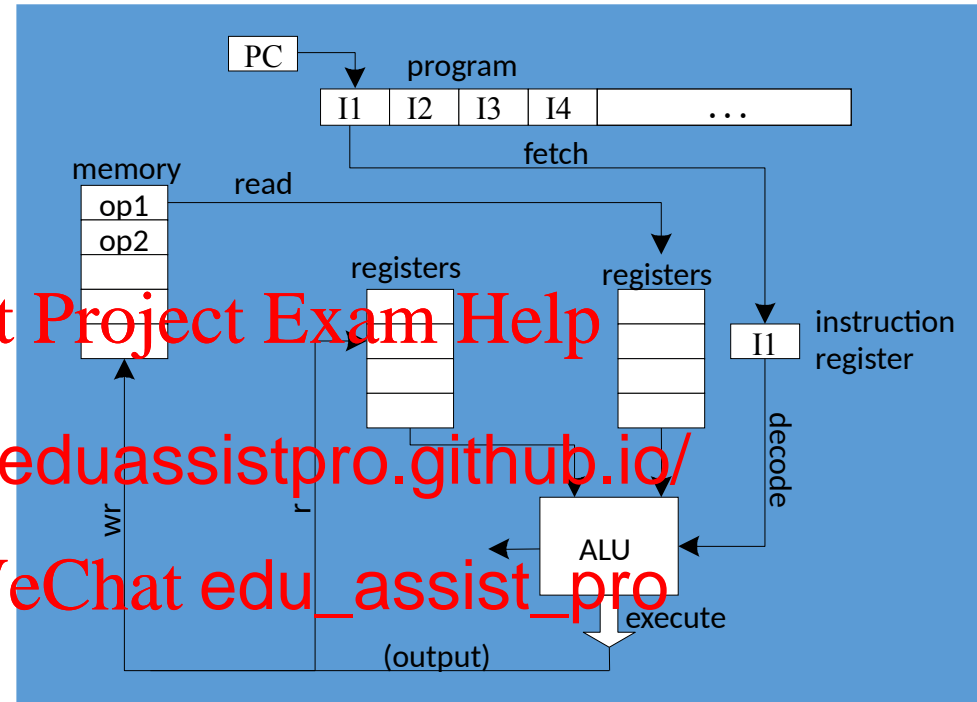
- Each machine language instruction is first fetched from the memory and stored in an **Instruction Register (IR)**.
- The address of the instruction is stored in a register called **Program Counter**. In computers this register is called the **Instruction Pointer** or **IP**.
- After the instruction is fetched, the **Program Counter** is incremented to point to the address of the next instruction.
- The fetched instruction is decoded (to determine what needs to be done) and executed by the CPU.

Instruction Execute Cycle



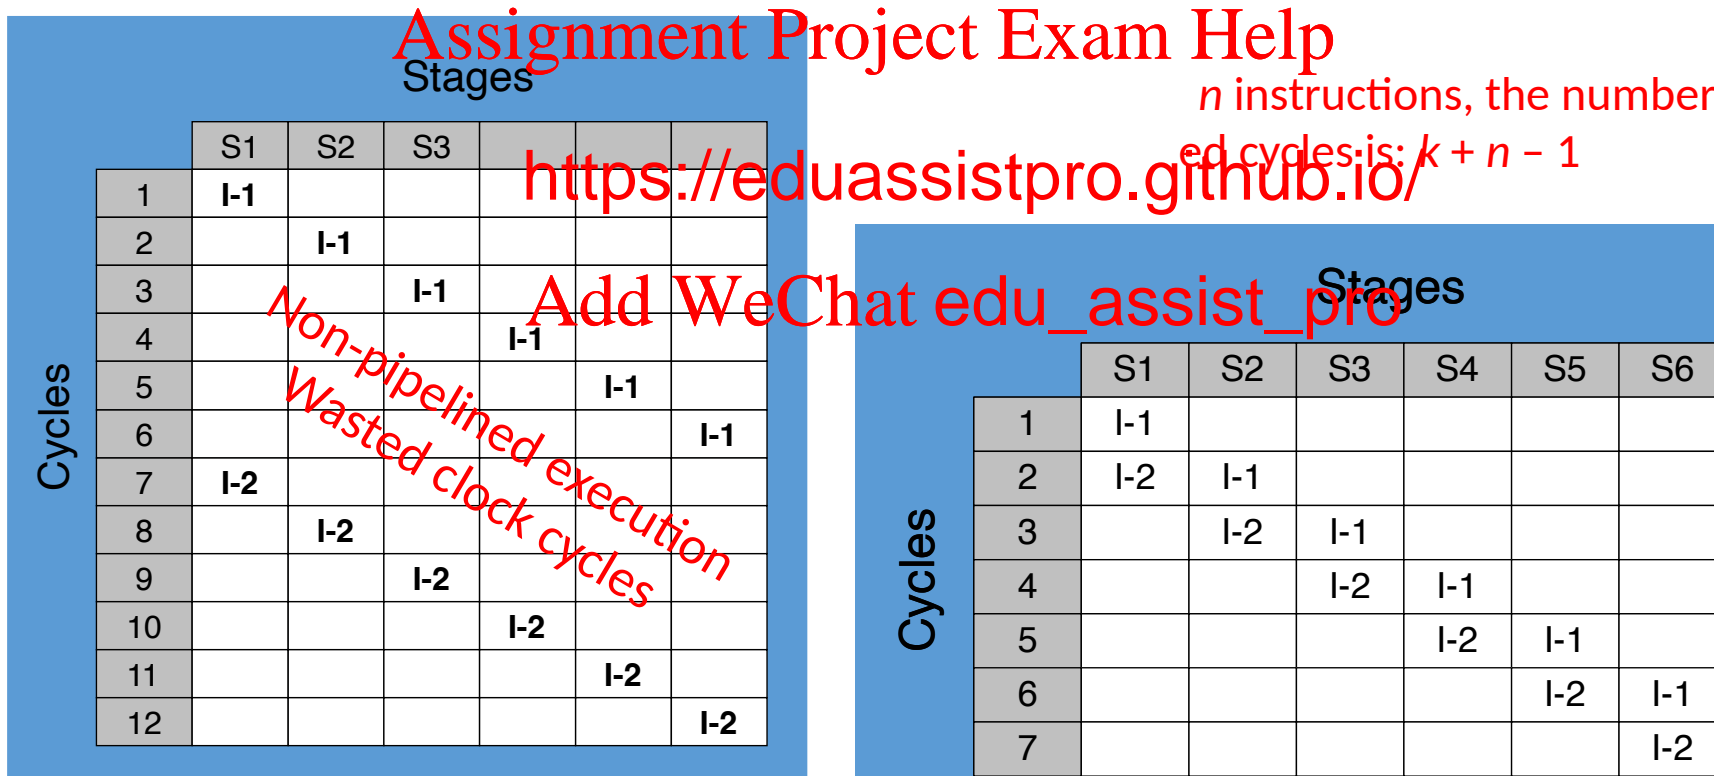
Instruction Execution Cycle – cont'd

- Instruction Fetch
- Instruction Decode
- Operand Fetch
- Execute
- Result Writeb



Pipelined Execution

- Instruction execution can be divided into stages
- Pipelining makes it possible to start an instruction before completing the execution of previous one



Wasted Cycles (pipelined)

- When one of the stages requires two or more clock cycles to complete, clock cycles are again wasted

- Assume that stage S4 is the execute stage
- Assume also that stage S4 requires 2 clock cycles to complete
- As more instructions enter the pipeline, wasted cycles occur
- For k stages, where one stage requires 2 cycles, n instructions require $k + 2n - 1$ cycles

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Cycle	Stages					
		S2	S3	S4	S5	S6
1						
2		-1				
3		-2	I-1			
4		-3	I-2	I-1		
5			I-3	I-1		
6				I-2	I-1	
7				I-2		I-1
8				I-3	I-2	
9				I-3		I-2
10					I-3	
11						I-3

Superscalar Architecture

- A superscalar processor has multiple execution pipelines
- The Pentium processor has two execution pipelines
 - Called U and V pipes

- In the following, stage S4 has 2 pipeli

- Each pipeline still requires 2 cycles
- Second pipeline eliminates wasted cycles
- For k stages and n instructions, number of cycles = $k + n$

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

		Stages						
				S3	u	v	S5	S6
Cycles								
	3	I-3	I-2	I-1				
	4	I-4	I-3	I-2	I-1			
	5		I-4	I-3	I-1	I-2		
	6			I-4	I-3	I-2	I-1	
	7				I-3	I-4	I-2	I-1
	8					I-4	I-3	I-2
	9						I-4	I-3
	10							I-4

Next ...

- Intel Microprocessors

- IA-32 Registers

Assignment Project Exam Help

- Instruction Execution

- IA-32 Memory Manag

<https://eduassistpro.github.io/>
Add WeChat edu_assist_pro

Modes of Operation

- Real-Address mode (original mode provided by 8086)
 - Only 1 MB of memory can be addressed, from 0 to FFFFF (hex)
 - Programs can access any part of main memory
 - MS-DOS runs in real address mode
- Protected mode
 - Each program can access up to 4 GB of memory
 - The operating system assigns memory to each running program
 - Programs are prevented from accessing each other's memory
 - Native mode used by Windows NT, 2000, XP, and Linux
- Virtual 8086 mode
 - Processor runs in protected mode, and creates a virtual 8086 machine with 1 MB of address space for each running program

Memory Segmentation

- Memory segmentation is necessary since the 20-bits memory addresses cannot fit in the 16-bits CPU registers
- Since x86 registers are 16-bits wide, a memory segment is made of 2^{16} consecutive bytes (i.e. 64K bytes)
- Each segment has a number identifier (i.e. we have segments numbered from 0 to 64K)
- A memory location within a memory segment is referenced by specifying its offset from the start of the segment. Hence the first word in a segment has an offset of 0 while the last one has an offset of FFFFh
- To reference a memory location its logical address has to be specified. The logical address is written as:
 - Segment number:offset
- For example, A43F:3487h means offset 3487h within segment A43Fh.

Program Segments

- Machine language programs usually have 3 different parts stored in different memory segments:
 - **Instructions**: This is the code part and is stored in the code segment
 - **Data**: This is the data part which is manipulated by the code and is stored in the data segment
 - **Stack**: The stack is a special memory structure used by the CPU to implement procedure calls. It follows a Last-In-First-Out (LIFO) structure. This data structure is stored in the stack.
- The segment numbers for the code segment, the data segment, and the stack segment are stored in the segment registers **CS**, **DS**, and **SS**, respectively.
- Program segments do not need to occupy the whole 64K locations in a segment

Real Address Mode

- A program can access up to six segments at any time

- Code segment
- Stack segment
- Data segment
- Extra segment

- Each segment is 64 KB

- Logical address

- Segment = 16 bits
- Offset = 16 bits

- Linear (physical) address = 20 bits

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Logical to Linear Address Translation

Linear address = Segment \times 10 (hex) + Offset

Example:

segment = A1F0 (hex)
offset = 04C0 (hex)
logical address = A1F0:04C0 (hex)
what is the linear address?

Solution:

A1F0	0	(add 0 to segment in hex)
+	04C0	(offset in hex)
<hr/>		
A23C0		(20-bit linear address in hex)

Segment Overlap

- There is a lot of overlapping between segments in the main memory.
- A new segment starts every 10h locations (locations). <https://eduassistpro.github.io/>
- Starting address of a segment always has a 0h LSD. [Add WeChat edu_assist_pro](#)
- Due to segments overlapping logical addresses are not unique .

Your turn . . .

What linear address corresponds to logical address 028F:0030?

Solution: $028F0 + 0030 = 02920$ (hex)

Always use <https://eduassistpro.github.io/>

What logical address corresponds to linear address 28F30h?

Many different **segment:offset** (logical) addresses can produce the same linear address 28F30h. Examples:

28F3:0000, 28F2:0010, 28F0:0030, 28B0:0430, . . .

Flat Memory Model

- Modern operating systems turn segmentation off
- Each program uses one 32-bit linear address space
 - Up to $2^{32} = 4$ GB of memory
 - Segment registers are defined
 - All segments are mapped to the same linear address space
- In assembly language, we use `.MODEL flat`
 - To indicate the Flat memory model
- A linear address is also called a virtual address
 - Operating system maps virtual address onto physical addresses
 - Using a technique called paging

Programmer View of Flat Memory

- Same base address for all segments

- All segments are mapped to the **same linear address space**

- EIP Register

- Points at next i

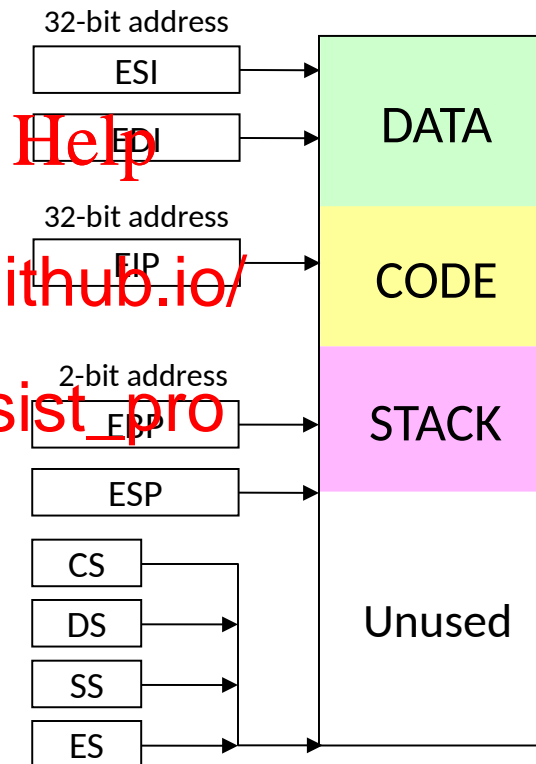
- ESI and EDI Reg

- Contain data addresses
 - Used also to index arrays

- ESP and EBP Registers

- ESP points at top of stack
 - EBP is used to address parameters and variables on the stack


Linear address space of a program (up to 4 GB)



base address = 0 for all segments

Protected Mode Architecture

- **Logical address** consists of
 - 16-bit **segment selector** (CS, SS, DS, ES, FS, GS)
 - 32-bit offset (EIP, ESP, EBP, ESI, EDI, EAX, EBX, ECX, EDX)
- Segment unit translates **logical address** to **linear address**
 - Using a **segment selector** and **base address** to calculate the **linear address**
 - Linear address is 32 bits (called **virtual address**)
- Paging unit translates **linear address** to **physical address**



Logical to Linear Address Translation

Upper 13 bits of segment selector are used to index the descriptor table

Assignment Project Exam Help

GDTR, LDTR

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

TI = Table Indicator

Select the descriptor table

0 = Global Descriptor Table

1 = Local Descriptor Table

Segment Descriptor Tables

- **Global descriptor table (GDT)**

- Only one GDT table is provided by the operating system
- GDT table contains segment descriptors for all programs
- Also used by the operating system
- Table is initialized during boot
- GDT table address is stored in the GDTR register
- Modern operating systems (Windows-XP) use one GDT table

- **Local descriptor table (LDT)**

- Another choice is to have a unique LDT table for each program
- LDT table contains segment descriptors for only one program
- LDT table address is stored in the LDTR register

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat: edu_assist_pro

Segment Descriptor Details

- **Base Address**

- 32-bit number that defines the starting location of the segment
- 32-bit Base Address + 32-bit Offset = 32-bit Linear Address

- **Segment Limit**

- 20-bit number that specifies the size of the segment
- The size is specified either in bytes or in 4 KB pages
- Using 4 KB pages, segment size can range from 4 KB to 4 GB

- **Access Rights**

- Whether the segment contains code or data
- Whether the data can be read-only or read & written
- Privilege level of the segment to protect its access

Segment Visible and Invisible Parts

- Visible part = 16-bit Segment Register
 - CS, SS, DS, ES, FS, and GS are visible to the programmer
- Invisible Part = Segment Descriptor (64 bits)
 - Automatically loaded from the descriptor table

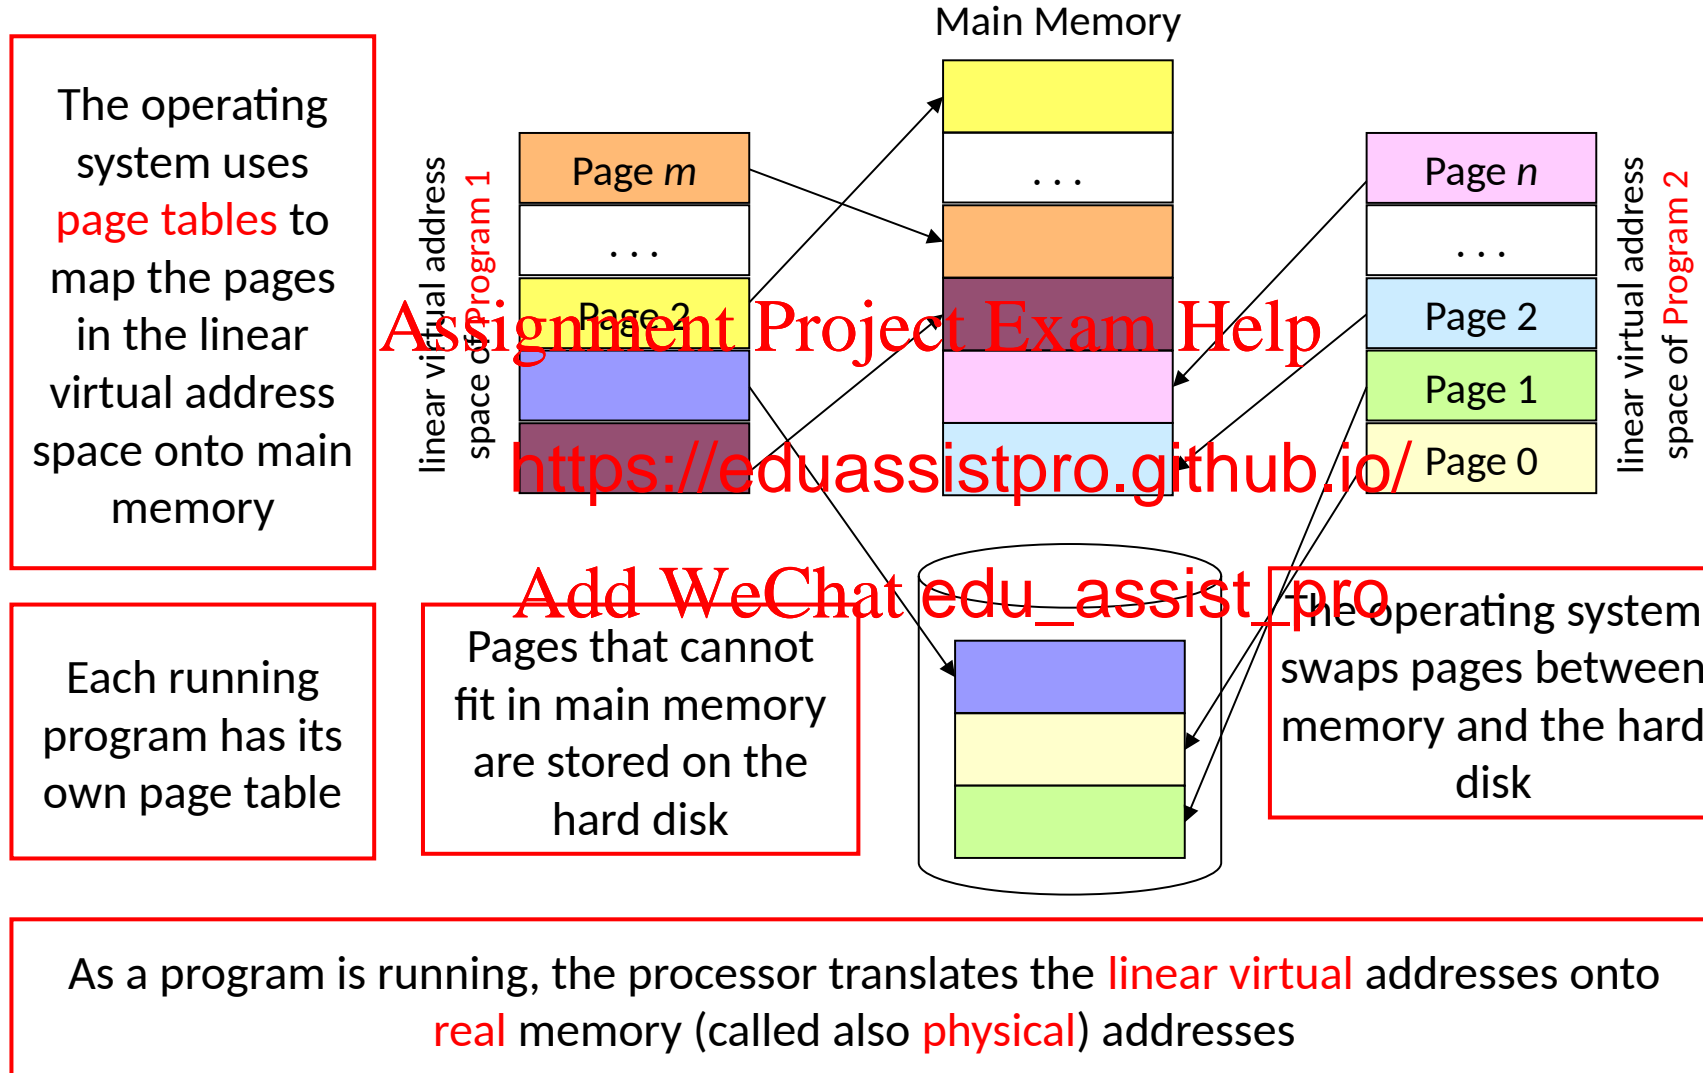
<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Paging

- Paging divides the linear address space into ...
 - Fixed-sized blocks called **pages**, Intel IA-32 uses 4 KB pages
- Operating system allocates main memory for pages
 - Pages can be spread all over main memory
 - Pages in main memory are shared by different programs
 - If main memory is full, pages are moved to the hard disk
- OS has a **Virtual Memory Manager**
 - Uses **page tables** to map the physical address of the running program
 - Manages the loading and unloading of pages
- As a program is running, CPU does address translation
- Page fault: issued by CPU when page is not in memory

Paging - cont'd



Components of an IA-32 Microcomputer

Assignment Project Exam Help

- <https://eduassistpro.github.io/>
- Video output
- Add WeChat edu_assist_pro
- Memory
- Input-output ports

Motherboard

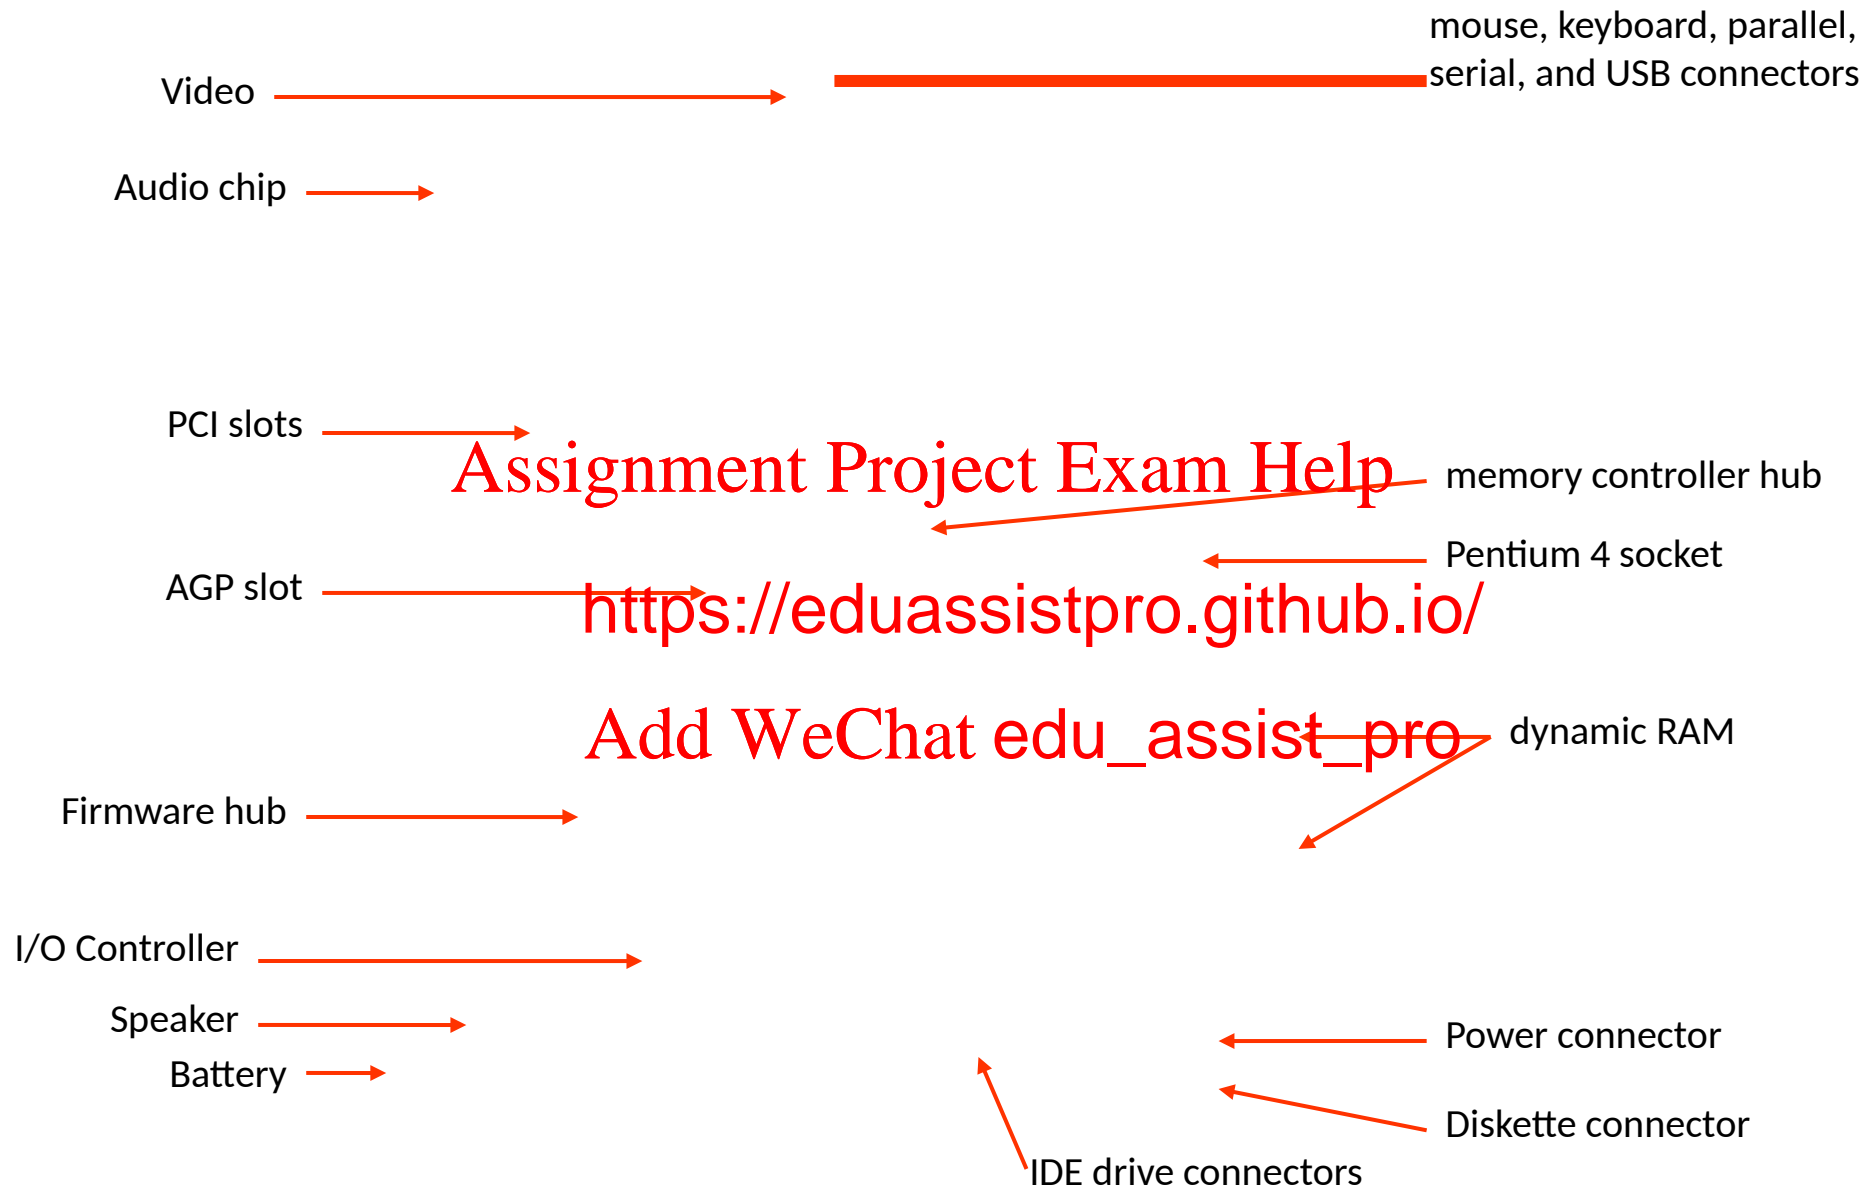
- CPU socket
- External cache memory slots
- Main memory slots
- BIOS chips
- Sound synthesizer chip (optional)
- Video controller chip (optional)
- IDE, parallel, serial, USB, video, keyboard, joystick, network, and mouse connectors
- PCI bus connectors (expansion cards)

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Intel D850MD Motherboard



Intel 965 Express Chipset

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Video Output

- Video controller
 - on motherboard, or on expansion card
 - AGP ([accelerated graphics port technology](#))*
- Video memory (VRAM)
- Video CRT
 - uses raster
 - horizontal retrace
 - vertical retrace
- Direct digital LCD monitors
 - no raster scanning required

* This link may change over time.

Sample Video Controller (ATI Corp.)

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Memory

- ROM
 - read-only memory
- EPROM
 - erasable programmable read-only memory
- Dynamic RAM (DRAM)
 - inexpensive; <https://eduassistpro.github.io/>
- Static RAM (SRAM)
 - expensive; used for cache memory
- Video RAM (VRAM)
 - dual ported; optimized for constant video refresh
- CMOS RAM
 - complimentary metal-oxide semiconductor
 - system setup information
- See: [Intel platform memory](#) (Intel technology brief: link address may change)

Input-Output Ports

- USB (universal serial bus)
 - intelligent high-speed connection to devices
 - up to 12 megabits/second
 - USB hub connects multiple devices
 - *enumerates* devices
 - supports <https://eduassistpro.github.io/>
- Parallel [Add WeChat edu_assist_pro](#)
 - short cable, high speed
 - common for printers
 - bidirectional, parallel data transfer
 - Intel 8255 controller chip

Input-Output Ports (cont)

- Serial
 - RS-232 serial port
 - one bit at a time
 - uses long cables and modems
 - 16550 U us receiver
transmitt
 - programmable in assem

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Next..

- General Concepts
 - IA-32 Processor Architecture
 - IA-32 Memory Management
 - Computer
 - Input
- <https://eduassistpro.github.io/>
Add WeChat edu_assist_pro

Levels of Input-Output

- Level 3: High-level language function
 - examples: C++, Java
 - portable, convenient, not always the fastest
- Level 2: Operating system
 - Application Programming Interface (API)
 - extended capabilities, lots of details to master
- Level 1: BIOS
 - drivers that communicate directly with devices
 - OS security may prevent application-level code from working at this level

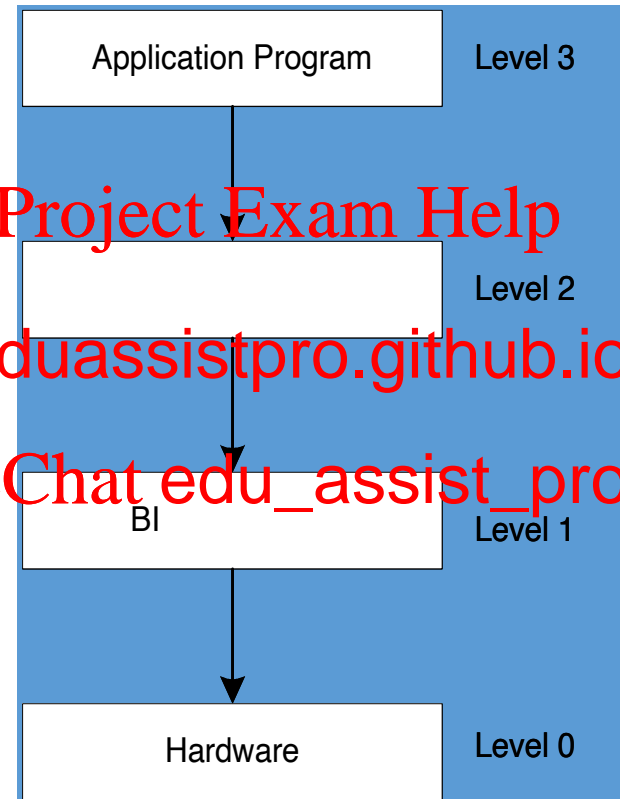
Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat: edu_assist_pro

Displaying a String of Characters

When a HLL program displays a string of characters, the following steps take place:



Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Programming levels

Assembly language programs can perform input-output at each of the following levels:

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Summary

- Central Processing Unit (CPU)
- Arithmetic Logic Unit (ALU)
- Instruction execution cycle
- Multitasking
- Floating Point Unit (FPU)
- Complex Instruction Set
- Real mode and Protected mode
- Motherboard components
- Memory types
- Input/Output and access levels

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro